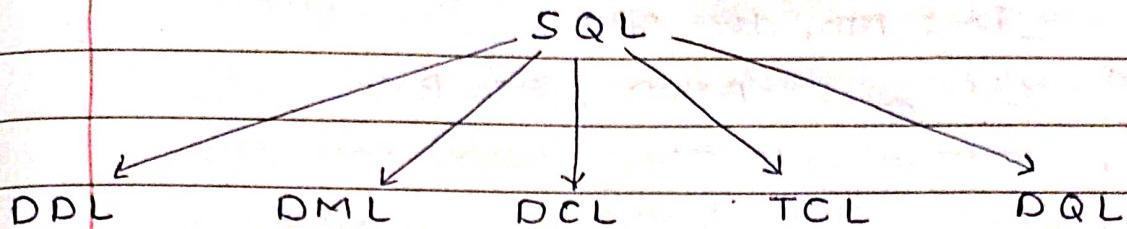


SQL (Structured Query Language)



CREATE INSERT GRANT COMMIT SELECT

ALTER UPDATE REVOKE ROLLBACK

DROP DELETE SAVE POINT

TRUNCATE

(Above are just basic commands)

change modification control Manage

Schema modification control Manage
OR
change transaction access of transactions
table into database data. of database

STRUCTURE SCHEMA & TABLES & INDEXES

OPERATORS

Arithmetic Comparison Logical Bitwise STRING

+,-,*,/,% =,>,<,<>,!= AND &,|,^, ||,+
OR

<=,>= NOT ~, <<, >>

(Concatenation)

• OTHER OPERATORS:

IN - checks within a set of values.

BETWEEN - checks within a range of values.

LIKE - Search for specified pattern.

% : 0,1 or multiple characters

_ : single character

\ - escape character

Basics

precedence

DOMS

Page No.

Date

/ /

Example:

%ANN% : Janny, Flannned

-is : his, dis, 8is

escape
character

-\% - : A%A, A%B

Precedence in SQL

FROM > JOIN > ON > USING > WHERE

> GROUP BY > HAVING > SELECT

> DISTINCT > WINDOW FUNCTIONS >

> ORDER BY > LIMIT/OFFSET

> FETCH FIRST/NEXT > OFFSET

Clauses [not common but important]

① DISTINCT : remove duplicates from result set.

② LIMIT : return maximum no. of rows from result set. (N) - parameter

③ COUNT (column-name) : count no. of rows.

* means count in column.

also count

1] doesn't difference any
2] null values.

Etc.] column, same as *

advance.

CASE

DOMS	Page No.
Date	/ /

④ CASE + WHEN + THEN + ELSE :

- switch and if/else of SQL.
- allows you to create expression, return different value based on evaluation of conditions.

Where to use: SELECT - return columns/value
WHERE - apply condition.

ORDER BY - sort result by cond.

HAVING - apply condition.

Example:

① SELECT A, B,
CASE COLC
WHEN 'HR' THEN 'HUMAN'
WHEN 'IT' THEN 'TECH'
ELSE 'OTHER'
END AS COL-C-CASE
FROM TableA ;

② Select id, name, department FROM employees
WHERE CASE
WHEN department = 'HR' THEN
WHEN " " = 'IT' THEN id < 5
ELSE TRUE
id < 3
END ;

③ ORDER BY CASE

WHEN department = 'HR' THEN 1
WHEN " " = 'IT' THEN 2
ELSE 3
END .

5. Two ways to write SQL Table Queries.

using , (Comma)

```
SELECT id, name, d.dname
```

```
FROM employees, department d
```

```
where id = d.did;
```

using JOIN

```
SELECT e.id, e.name, d.dname
```

```
FROM employee e
```

```
JOIN department d ON e.id = d.id;
```

I prefer JOIN method as it is more intuitive, good and easy to grasp.

6. UNION & UNION ALL

not include also include
duplicate duplicates.

Both used to combine two result sets.

conditions: → Data type of column must match

→ No. of columns must match.

advance
coalesce
aggregation functions

DOMS	Page No.
Date	/ /

⑦ COALESCE / NVL

NVL is specific to oracle and takes

return the first non-null value from a list of expressions. only two arguments - null value from a list of expressions. NVL(expr, replace-val)

COALESCE (Ex1, Ex2, Ex3...ExN)

eg:

```
SELECT id,  
COALESCE(middle_name, first_name, last_name) AS name  
FROM employees;
```

mostly used with

Aggregate fun. Group By Clause

perform calculation over multiple rows &

return a single value.

COUNT() SUM() AVG() MAX() VARIANCE()

MIN() STDDEV()

- applied over columns.
- with GROUP BY clause to calculate value according to groups.
- with HAVING / WHERE clause to filter out values.

Joins

- they are used to combine two or more tables or result sets.

1. JOIN / INNER JOIN

return record having matching value in both tables.

2. LEFT OUTER RIGHT OUTER FULL OUTER

INNER JOIN

+

all records

from left
table,

non match

INNER JOIN

+

all records

from right
table,

non match

Value

replace by

NULL

INNER JOIN

+

all records

from left
table

non match

all records

from right
table

table

3. CROSS JOIN

return cartesian product of two tables

i.e., combine each row of left table with each row of right table

4. NATURAL JOIN

is to INNER JOIN but ON condition is

determined by SQL itself based on column names, primary keys etc.

NOT recommended, inconsistent results.

5. SELF JOIN

- INNER JOIN but it is join with itself i.e. same table.
- Use case: When you have to link one records of table with another records of same table.

Key words

1. JOIN ON
2. LEFT JOIN / RIGHT / FULL
3. CROSS JOIN ON
4. NATURAL JOIN ON
5. SELF JOIN ON

Sub Queries

1.

Scalar

return single

call stored

Multiple

either multiple

rows or columns

or single column

multiple rows.

- Both are inter dependent.
- not related to outer sub query.
- use () in where clause to compare multiple checks or conditions.

2.

Correlated

a subquery which is related to outer subquery.

for each outer query record, inner sub query is executed. like nested loops.

you can nest as many as queries you want but keep in mind it is computationally expensive.

- All subqueries can be used anywhere like in WHERE, JOIN, HAVING, SELECT.

(imp)

in SELECT clause, you must do only scalar subqueries otherwise will throw error.

OR basically, correlated query throw error in SELECT.

WITH clause / CTE (common table expressions)

one time

- It works like ^{variables} in other languages like CPP, python etc.
- It runs a query and stores its result in RAM until the execution of that particular query.
- Returns the result set.
- Increase readability and speed.
- Time span : current Query Execution only.
- Very helpful in recursive queries.

syntax:

WITH CTE_NAME AS (

query

) main query.

Eg: WITH employeedepartments AS (

SELECT e.id, e.name

FROM employee AS e

JOIN department d ON e.id = d.id

)

SELECT * FROM employeedepartments;

- CTE are also used to write recursive queries.

aduanc
recursion.

DOMS Page No.
Date / /

Recursive Queries

- Syntax:

WITH RECURSIVE cte-name AS (

-- Base/Anchor member

[
 Query

UNION OR UNION ALL

-- Recursive member

[
 Query (inside FROM there must be
 cte-name as it reference
 same result set)

)

main query.

- Anchor member
 - provide base result and starting point

- Recursive member

→ references CTE itself & iteratively adds rows to result set.

→ it also have termination condition, or query will end when no result is found.

Eg:

WITH RECURSIVE eh AS (

-- Anchor

SELECT e.id, e.name, e.mid, 1 AS level

FROM employees

WHERE manager_id IS NULL

UNION ALL

-- Recursive part of query

-- Recursive part of query

SELECT e.id, e.name, e.mid, level+1

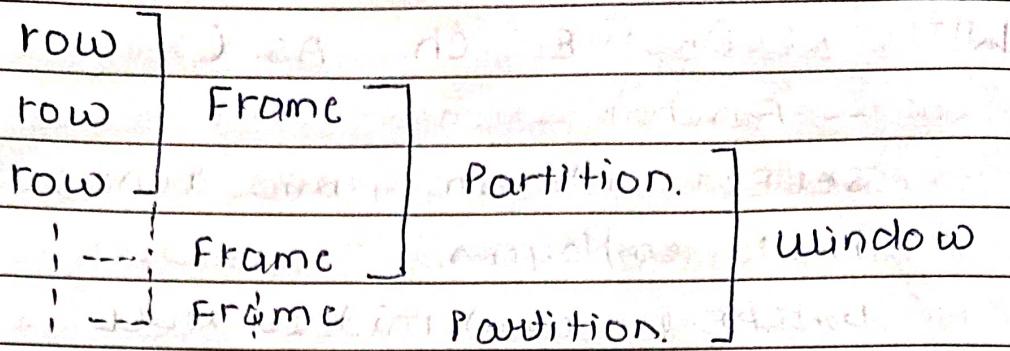
FROM employees e

JOIN eh AS eh1 ON e.mid = eh1.id

SELECT e.id, e.name, e.mid, level

FROM eh

Window Functions.



- running rows makes a frame

1 row = 1 frame

2 row = frame 2

3 row = frame 3 etc

or it can be depend on condition.

but the point is 0 — i^{th} current

0 — $(i+1)^{th}$ next frame

- window func are used to perform calculations across a set of table (window) that are related to current row

OVER PARTITION BY ORDER BY

↓ ↓ ↓
Defines the like a group by order rows
window to for window over within
operation. which other func partition.
can be applied.

1. RANK()

→ Assign rank to each row in partition.
→ Assign same rank to duplicates but skip next values.

Eg:

value: 10 20 20 43

RANK: 1 2 2 3

2. DENSE_RANK()

→ same as Rank fun but do not skip values. It goes like 1 2 2 3

3. ROW_NUMBER()

→ assign each row a unique row number

4. Aggregate funs can also be used using OVER clause to apply them on each partition.

5. LAG()

- access data from previous rows - access data from next rows.

- LAG(column, N) LEAD(column, N)

6. FIRST_VALUE(column)

LAST_VALUE(column) } respective value

NTH_VALUE(column, N) } from each partition.

7 NTILE(N): SQL try to divide in equal partition.

8. CUME_DIST() PERCENT_RANK()

↓ gives the cumulative distribution rank of row within of row without partition i.e. of rows. How many rows i.e. How much. How much or higher distribution current, relative to current row till now holds. row.

FORMULA: $\frac{(\text{No. of rows with value} \leq \text{current Row No. - 1})}{(\text{total - 1})}$

Range: $0 \leq \text{value} \leq 1$ $0 \leq \text{value} \leq 1$

NOTE: In case of duplicate records, row of last record is considered.

9. FRAME clause

- FRAME word is not used
- used at end of OVER clause as last arguments.

Syntax:

(A) BETWEEN (B) PRECEDING AND (C)
Following.

window fn)

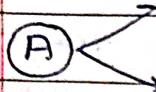
frame clause

DOMS

Page No.

Date / /

RANGE: take account of duplicates.

(A)  eg: last 2 : [5 4 4] 3 2

ROWS: specific to rows only.

eg: last 2 : [5 4] 4 3 2

(B) → unbounded : from starting of partition.

→ n : from last n of partition.

* (range) rows

(there are many more)

unbounded : till end of current partition.

n : next n (range / rows) of current partition.

current rows: till current row

excluding and after (omit word following)

WHY?

Because in fn) like last_value()

frame considers the running rows.

	C1	C2	Frame	Output	we want
partition	A	1	1	1	4
	B	2	2	2	4
	C	3	3	3	4
	D	4	4	4	4

F1 : A $\therefore O = 1$

F2 : AB $\therefore O = 2$

F3 : ABC

F4 : ABCD

as those funs are applied frame by frame. even though they are applied over partition, as per saying, but executed frame by frame.

By default FRAME clause of

LAST_VALUE is

- * Range between unbounded preceding and current row.
 - ∴ it will generate wrong results.
- Hence, we need frame clause.

Examples:

1. Select id, name, salary,
 RANK() OVER (ORDER BY salary desc)
 FROM employees;

2. Select id, date, amount,
 AVG(amount) OVER (
 order by date
 ROWS between 2 preceding and
 current row
) as moving_avg
 FROM orders;

3. Select id, salary,
 PERCENT_RANK() OVER() AS PR
 FROM employees.

window fun
window keyword

DOMS	Page No.
Date	/ /

NOTE: if you have common window conditions i.e. inside over brackets. then you can omit common conditions and write it in window keyword fun

Example:

At end of query mention. like

1. At end of query mention OVER (PARTITION BY)

2. WINDOW W AS (conditions)

URC.

So, this is the order of the window function

Query.

Select * ,

FIRST_VALUE (Pn) OVER (W) AS FV,

LAST_VALUE (Pn) OVER (W) AS LV

From products.

where PC = 'Phonex'

WINDOW W AS (partition by PC

order by price desc

range between 2 preceding

and 2 following);

crosstab # pivot
rows to columns

DOMS

Page No.

Date

/

/

CROSSTAB / PIVOT in other DB



used to transform rows into columns.

In crosstab must add Unc to create extension.

Unc:

CREATE Extension IF NOT Exists tablefunc

syntax:

Select *

FROM crosstab (

'SELECT [row identifier], *

[column to pivot], *

important

[data column] *

all 3

FROM [table]

ORDER BY 1,2

) AS ct(

[row identifier], [data type],

[value1] [data type],

[value2] [data type],

);

focus to columns

DOMS	Page No.
Date	/ /

Example:

CREATE EXTENSION IF NOT EXISTS tablefunc;

SELECT * FROM crossTab(

'SELECT salesperson, month, sales
FROM sales_data
ORDER BY 1, 2'

) AS ct (salesperson int, Jan int,
Feb int, Mar int);

→ this line works as magic

why?

because postgres just pick values and
place it according to new columns Hence,
we have to sort data accordingly.

Views

Views are like CTE's but CTE associated to only single query. & SQL don't store schema of CTE as it is used in one query only.

they are like virtual tables.

- schema is stored
- permanent object
- always present, until explicitly dropped
- will run every time, it is used
it uses user stored variable but evaluated every time it is used.

Basic DDL syntax:

* CREATE VIEW v-name
AS (Query)

1. CREATE OR REPLACE

if already present then

- can't change column name, datatype,
- add new column only at end. order.

2. ALTER VIEW

change structure

3. DROP VIEW

delete view

NOTE: (if table structure changed view needs to be updated as it stores previous version of table)

Updatable views

Concept

DOMS	Page No.
Date	/ /

Updatable Views:

only views which

1. Created by only one table can be updated.
2. NOT have DISTINCT, Aggregate functions
3. NOT have GROUP BY
4. NOT have CTE/ WITH clause.
5. NOT have JOINS
6. NOT have subqueries.

Eg: Update VIEW_A

SET d = "Market"

WHERE id = 1;

- changes reflected back into table also.
- you can also add check option so that where condition must met before updation.

Eg:

Query

WHERE condition.

WITH CHECK OPTION;

materialized views
optimization.

Materialized View

use word MATERIALIZED before view
saves both schema and result set.
Hence,
performance very fast
but takes extra space.
→ we can manually refresh them
if new data is added to table.

CMD:

Refresh Materialized View MVA;

Syntax:

CREATE MATERIALIZED VIEW MVA AS

Query

A. Trigger

In SQL, a trigger is a special type of stored procedure that automatically executes or "fires" in response to specific events or changes in a database. Triggers are used to enforce business rules, automate tasks, and maintain data integrity.

Features

1. **Automatic Execution:**
Triggers are executed automatically in response to events like `INSERT`, `UPDATE`, or `DELETE` on a table.
2. **Event-Driven:**
They are triggered by specific database events and can be set to run before or after the event occurs.
3. **Maintain Data Integrity:**
Triggers can enforce rules and constraints to ensure data consistency and integrity.
4. **Complex Logic:**
They allow for the execution of complex logic that cannot be easily handled with constraints alone.

Types

1. **BEFORE** Triggers:
Executed before an `INSERT`, `UPDATE`, or `DELETE` operation on a table.
2. **AFTER** Triggers:
Executed after an `INSERT`, `UPDATE`, or `DELETE` operation on a table.
3. **INSTEAD OF** Triggers:
Executed in place of the triggering operation, allowing for custom logic to replace the standard operation.

Syntax

```
CREATE FUNCTION trigger_function()
RETURNS TRIGGER
AS $$
BEGIN
    -- trigger logic here
    RETURN NEW;
END;
$$ LANGUAGE plpgsql;
```

```
CREATE TRIGGER trigger_name
BEFORE INSERT ON table_name
FOR EACH ROW
EXECUTE FUNCTION trigger_function();
```

Example

```
CREATE OR REPLACE FUNCTION log_salary_changes()
RETURNS TRIGGER AS $$
BEGIN
    -- Insert a record into salary_log for each update
    INSERT INTO salary_log (employee_id, old_salary, new_salary)
    VALUES (OLD.employee_id, OLD.salary, NEW.salary);

    -- Return the new row
    RETURN NEW;
END;
$$ LANGUAGE plpgsql;
```

```
CREATE TRIGGER trigger_salary_changes
AFTER UPDATE OF salary
ON employees
FOR EACH ROW
EXECUTE FUNCTION log_salary_changes();
```

B. Stored Procedures

A stored procedure is a set of SQL statements that perform a specific task, which is stored in the database and can be executed by calling the procedure. Unlike functions, stored procedures do not return a value but can perform operations such as modifying data or managing transactions.

Syntax

```
CREATE PROCEDURE procedure_name (param1 type1, param2 type2, ...)
LANGUAGE plpgsql
AS $$
```

```
DECLARE
  Var1 DATATYPE;
  var2 DATATYPE;
  ...
  var_n DATATYPE;
```

```
BEGIN
  -- Procedure logic here
END;
$$;
```

Example

```
CREATE PROCEDURE update_department_average(department_name VARCHAR)
LANGUAGE plpgsql
AS $$
```

```
DECLARE
  total_salary DECIMAL;
  employee_count INT;
  avg_salary DECIMAL;

BEGIN
  -- Calculate total salary and employee count for the department
  SELECT SUM(salary), COUNT(*)
  INTO total_salary, employee_count
  FROM employees
  WHERE department = department_name;

  -- Calculate average salary
  IF employee_count > 0 THEN
    avg_salary := total_salary / employee_count;
  ELSE
    avg_salary := 0;
  END IF;
```

```

-- Update department table with average salary (assuming there's a department table)
UPDATE departments
SET average_salary = avg_salary
WHERE department_name = department_name;
END;
$$;

-- Calling the procedure
CALL update_department_average('Engineering');

```

Parameters

1. IN Parameters: Used to pass values into the procedure.
2. OUT Parameters: Used to return values from the procedure to other procedure.
3. INOUT Parameters: Can be used to pass values in and return values out.

C. Functions

A function in SQL is a stored routine that returns a value. Functions can be used in SQL statements and can perform computations, return results, and be used in SELECT, WHERE, or ORDER BY clauses.

Syntax

```

CREATE FUNCTION function_name (param1 type1, param2 type2, ...)
RETURNS return_type
LANGUAGE plpgsql
AS $$
```

```

DECLARE
  var1 DATATYPE;
  var2 DATATYPE;
  ...
  var_n DATATYPE;
```

```

BEGIN
  -- Function logic here
  RETURN return_value;
END;
$$;
```

Example

```

CREATE FUNCTION get_total_salary(department_name VARCHAR)
RETURNS DECIMAL
LANGUAGE plpgsql
AS $$
```

```

DECLARE
  total_salary DECIMAL;
BEGIN
  -- Calculate total salary for the department
  SELECT SUM(salary)
  INTO total_salary
  FROM employees
  WHERE department = department_name;
```

```
-- Return the total salary
```

```
RETURN total_salary;
END;
$$;

-- Calling the function
SELECT get_total_salary('Engineering');
```

Difference between Function and Procedure

1. Return Value:

Function: Returns a single value or a set of values (e.g., scalar, table).

Procedure: Does not return a value directly; can use `OUT` parameters to return values.

2. Usage in SQL Statements:

Function: Can be used within SQL statements (e.g., `SELECT`, `WHERE`).

Procedure: Cannot be used within SQL statements; called using `CALL`.

3. Purpose:

Function: Typically used for computations and data retrieval.

Procedure: Used for performing operations, including data modification and managing transactions.

4. Transactional Control:

Function: Limited control over transactions (cannot commit or roll back transactions within the function).

Procedure: Can include explicit transaction control (e.g., `COMMIT`, `ROLLBACK`).

5. Side Effects:

Function: Ideally should not have side effects (e.g., modifying database state).

Procedure: Often used to perform operations that change the database state.